

WK8_HW8

Anon Skywalker

10/13/2020

Question 11.1 Stepwise, Lasso, and Elastic Net Variable Selection

Base Model

To begin, I start with our most default regression model using the uscrime data with all predictors. This will serve as a base model for which we can begin to compare our future models with that will utilize different variable selection tactics.

Base Model Stats:

1. All Predictors
2. AIC of 650
3. R-squared of: 0.353

```
set.seed(1)
lm.crime.default <- lm(Crime~., data = uscrime)
#summary(lm.crime.default)

# cross validate using 4 folds
cv.lm.default <- cv.lm(lm.crime.default, data = uscrime, m = 4, plotit = FALSE)
```

```

## Analysis of Variance Table
##
## Response: Crime
##      Df Sum Sq Mean Sq F value Pr(>F)
## M      1   55084   55084    1.26  0.2702
## So      1   15370   15370    0.35  0.5575
## Ed      1  905668  905668   20.72 7.7e-05 ***
## Po1     1 3076033 3076033   70.38 1.8e-09 ***
## Po2     1  153024   153024    3.50  0.0708 .
## LF      1   61134   61134    1.40  0.2459
## M.F     1  111000  111000    2.54  0.1212
## Pop     1   42649   42649    0.98  0.3309
## NW      1   14197   14197    0.32  0.5728
## U1      1    7065    7065    0.16  0.6904
## U2      1  269663  269663    6.17  0.0186 *
## Wealth  1   34748   34748    0.79  0.3795
## Ineq    1  547423  547423   12.52 0.0013 **
## Prob    1  222620  222620    5.09  0.0312 *
## Time    1   10304   10304    0.24  0.6307
## Residuals 31 1354946  43708
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##
## fold 1
## Observations in test set: 11
##      2  9  14  16  20  22  26  38  41  44  47
## Predicted 1474 689 780 1006 1227.8 657 1977.4 562.7 824 1121 992
## cvpred    1535 706 867 1100 1298.9 931 2043.3 602.8 757 1257 1159
## Crime      1635 856 664 946 1225.0 439 1993.0 566.0 880 1030 849
## CV residual 100 150 -203 -154 -73.9 -492 -50.3 -36.8 123 -227 -310
##
## Sum of squares = 512057    Mean square = 46551    n = 11
##
## fold 2
## Observations in test set: 12
##      1  3  6  11  19  25  28  29  30  33  35  39
## Predicted 755.0 322 793 1161 1146 605.9 1258.48 1287 703 841 738 839.3
## cvpred    727.7 265 920 1082 1449 535.1 1219.78 1534 634 784 886 868.7
## Crime      791.0 578 682 1674 750 523.0 1216.00 1043 696 1072 653 826.0
## CV residual 63.3 313 -238 592 -699 -12.1 -3.78 -491 62 288 -233 -42.7
##
## Sum of squares = 1382466    Mean square = 115205    n = 12
##
## fold 3
## Observations in test set: 12
##      4  5  10 12  13  15 17  34  37  40  42  45
## Predicted 1791 1167 736.5 722 733 903 393 971.5 971 1131.5 326.3 617
## cvpred    1576 1021 745.1 824 912 1050 103 823.4 1392 1186.8 -85.5 848
## Crime      1969 1234 705.0 849 511 798 539 923.0 831 1151.0 542.0 455
## CV residual 393 213 -40.1 25 -401 -252 436 99.6 -561 -35.8 627.5 -393
##
## Sum of squares = 1491541    Mean square = 124295    n = 12

```

```
##
## fold 4
## Observations in test set: 12
##           7      8     18     21     23     24     27     31     32     36     43     46
## Predicted   934.2 1362   844 774.9  958 869 279.5 388.0   808 1138 1134   827
## cvpred      1055.1 1123 1189 725.3  922 851 272.7 433.1   953  852 1324   984
## Crime       963.0 1555   929 742.0 1216 968 342.0 373.0   754 1272   823   508
## CV residual -92.1  432 -260   16.7  294 117   69.3 -60.1 -199  420 -501 -476
##
## Sum of squares = 1065774    Mean square = 88814    n = 12
##
## Overall (Sum over all 12 folds)
##      ms
## 94720
```

```
#A function for calculating r-squared for each improved regression model
R2 <- function(ms) {

  sse <- ms*nrow(uscrime)
  sst<- sum((uscrime$Crime - mean(uscrime$Crime))^2)
  R2 <- 1 - sse/sst
  #0.615 R-SQUARED
  return(R2)

}

#calculate for r-squared for base model
R2(94720)
```

```
## [1] 0.353
```

```
#650
AIC(lm.crime.default)
```

```
## [1] 650
```

Stepwise begins here:

Now that we have a base model to compare with, we can dive into the different types of variable selection methods. Beginning with Stepwise regression, where predictive variables are either added or subtracted at each step of the selection process. To achieve stepwise efficiently, I utilized the CARET packages `train()` function, which sets up a grid of tuning parameters for a number of classification and regression routines, fits each model and calculates a resampling based performance measure. The `method= "leapSeq"` is the stepwise regression method. The `trainControl` is a parameter used to apply cross validation with 10 folds onto each model `train()` comes up with.

After running the stepwise approach, the resulting model seemed to be best fitted with 6 predictors as can be seen by the `$bestTune` provided by the model. The final model returned by the StepWise approach utilized the `M + Ed + Po1 + U2 + lneq + Prob` predictors of which I used in the `step.improv` model. Comparing this to our base model, it is a improvement in model quality that is notable, with a AIC of 10 points lower. Note: While the coefficients are not

the same via this approach, we standardize the method of using the stepwise/lasso/elastic net variable selection processes to take their chosen predictors and use them with `lm()` to be able to test their AIC and r-squared for better model comparison. This method will be applied to each model.

Stepwise Model Results:

1. Preds in improv model: M + Ed + Po1 + U2 + Ineq + Prob
2. AIC = 640
3. R-Squared = 0.671

```
set.seed(1)
# Set up repeated k-fold cross-validation
train.control <- trainControl(method = "cv", number = 10)
#"leapSeq", to fit linear regression with stepwise selection
step.model <- train(Crime~., data = uscrime, method = 'leapSeq',
                    tuneGrid = data.frame(nvmax = 1:8), trControl = train.control)

#r-Squared of 0.61 for 6 vars!
step.model$results
```

##	nvmax	RMSE	Rsquared	MAE	RMSESD	RsquaredSD	MAESD
## 1	1	271	0.584	226	110.5	0.337	96.7
## 2	2	254	0.606	202	84.8	0.281	64.1
## 3	3	239	0.703	186	88.7	0.101	59.2
## 4	4	276	0.664	216	104.7	0.251	79.0
## 5	5	252	0.565	207	80.6	0.280	62.7
## 6	6	233	0.785	185	102.3	0.148	81.6
## 7	7	240	0.729	191	95.4	0.182	74.0
## 8	8	259	0.699	203	112.0	0.203	90.3

```
#best # of variables
step.model$bestTune
```

```
##    nvmax
## 6      6
```

```
#best model appears to be with 6 predictors!
step.final <- step.model$finalModel
```

```
## = Significance. Model Chose: M + Ed  Po1  M.F U1  U2 Ineq  Prob
summary(step.final)
```

```
## Subset selection object
## 15 Variables (and intercept)
##      Forced in Forced out
## M      FALSE      FALSE
## So      FALSE      FALSE
## Ed      FALSE      FALSE
## Po1     FALSE      FALSE
## Po2     FALSE      FALSE
## LF      FALSE      FALSE
## M.F     FALSE      FALSE
## Pop     FALSE      FALSE
## NW      FALSE      FALSE
## U1      FALSE      FALSE
## U2      FALSE      FALSE
## Wealth  FALSE      FALSE
## Ineq    FALSE      FALSE
## Prob    FALSE      FALSE
## Time    FALSE      FALSE
## 1 subsets of each size up to 6
## Selection Algorithm: 'sequential replacement'
##      M  So  Ed  Po1 Po2 LF  M.F Pop NW  U1  U2  Wealth Ineq Prob Time
## 1 ( 1 ) " " " " " " "*" " " " " " " " " " " " " " " " " " " " "
## 2 ( 1 ) " " " " " " "*" " " " " " " " " " " " " " " " " " " " "
## 3 ( 1 ) " " " " "*" "*" " " " " " " " " " " " " " " " " " " " "
## 4 ( 1 ) "*" "*" "*" "*" " " " " " " " " " " " " " " " " " " " "
## 5 ( 1 ) "*" " " "*" "*" " " " " " " " " " " " " " " " " " " " "
## 6 ( 1 ) "*" " " "*" "*" " " " " " " " " " " " " " " "*" " " " " " "
```

```
#display final stepwise model coefficients
coef(step.final, 6)
```

```
## (Intercept)          M          Ed          Po1          U2          Ineq
##    -5040.5        105.0        196.5        115.0        89.4        67.7
##          Prob
##    -3801.8
```

```
#create LM model with newfound variable combination
step.improv <- lm(Crime~M + Ed + Po1 + U2 + Ineq + Prob, data = uscrime)
summary(step.improv)
```

```
##
## Call:
## lm(formula = Crime ~ M + Ed + Po1 + U2 + Ineq + Prob, data = uscrime)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -470.7  -78.4  -19.7   133.1   556.2
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -5040.5      899.8   -5.60  1.7e-06 ***
## M              105.0       33.3    3.15  0.0031 **
## Ed             196.5       44.8    4.39  8.1e-05 ***
## Po1            115.0       13.8    8.36  2.6e-10 ***
## U2              89.4       40.9    2.18  0.0348 *
## Ineq           67.7       13.9    4.85  1.9e-05 ***
## Prob        -3801.8     1528.1   -2.49  0.0171 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 201 on 40 degrees of freedom
## Multiple R-squared:  0.766, Adjusted R-squared:  0.731
## F-statistic: 21.8 on 6 and 40 DF, p-value: 3.42e-11
```

```
# ms of 48203
cv.lm.step <- cv.lm(step.improv, data = uscrime, m=4, plotit = FALSE)
```

```

## Analysis of Variance Table
##
## Response: Crime
##      Df Sum Sq Mean Sq F value Pr(>F)
## M      1  55084   55084    1.37 0.24914
## Ed      1 725967  725967   18.02 0.00013 ***
## Po1     1 3173852 3173852   78.80 5.3e-11 ***
## U2      1  217386   217386    5.40 0.02534 *
## Ineq    1  848273   848273   21.06 4.3e-05 ***
## Prob    1  249308   249308    6.19 0.01711 *
## Residuals 40 1611057   40276
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##
## fold 1
## Observations in test set: 11
##      2  9  14  16  20  22  26  38  41  44  47
## Predicted 1388 719 713.6 1004.4 1203.0 728 1789 544.4 796 1178 976
## cvpred    1355 731 731.1 1023.2 1187.6 771 1720 588.4 763 1150 970
## Crime     1635 856 664.0 946.0 1225.0 439 1993 566.0 880 1030 849
## CV residual 280 125 -67.1 -77.2 37.4 -332 273 -22.4 117 -120 -121
##
## Sum of squares = 334042    Mean square = 30367    n = 11
##
## fold 2
## Observations in test set: 12
##      1  3  6  11  19  25  28  29  30  33  35  39
## Predicted 810.8 386 730 1118 1221 579.1 1259.0 1495 668.0 874 808 786.7
## cvpred    716.9 296 888 1241 1363 504.3 1208.7 1711 614.2 792 919 736.6
## Crime     791.0 578 682 1674 750 523.0 1216.0 1043 696.0 1072 653 826.0
## CV residual 74.1 282 -206 433 -613 18.7 7.3 -668 81.8 280 -266 89.4
##
## Sum of squares = 1300449    Mean square = 108371    n = 12
##
## fold 3
## Observations in test set: 12
##      4  5  10  12  13  15  17  34  37  40  42  45
## Predicted 1897.2 1269.8 787.3 673 739 828 527.4 997.5 992 1140.8 369 622
## cvpred    1916.6 1282.8 791.8 680 778 867 483.3 998.2 1037 1190.7 317 656
## Crime     1969.0 1234.0 705.0 849 511 798 539.0 923.0 831 1151.0 542 455
## CV residual 52.4 -48.8 -86.8 169 -267 -69 55.7 -75.2 -206 -39.7 225 -201
##
## Sum of squares = 261503    Mean square = 21792    n = 12
##
## fold 4
## Observations in test set: 12
##      7  8  18  21  23  24  27  31  32  36  43  46
## Predicted 733 1354 800 783 938 919.4 312.2 440 774 1102 1017 748
## cvpred    708 1319 771 759 909 896.3 316.2 426 740 1093 1027 723
## Crime     963 1555 929 742 1216 968.0 342.0 373 754 1272 823 508
## CV residual 255 236 158 -17 307 71.7 25.8 -53 14 179 -204 -215
##

```

```
## Sum of squares = 369549    Mean square = 30796    n = 12
##
## Overall (Sum over all 12 folds)
##      ms
## 48203
```

```
#Calculate Rsquared for Stepwise model
R2(48203)
```

```
## [1] 0.671
```

```
#compare AICs between default and improv model: improv wins!
AIC(step.improv)
```

```
## [1] 640
```

Lasso Model

To begin with the lasso model, it is vital that the data is scaled due to the glmnet's need for matrix type data in the x and y parameters. This presented the challenge of not destroying the response column "Crime" and categorical column "So". To prevent this issue, I scaled all columns except for these two and then added them back in. Because of this, the parameter "standardize" was not used in my cv.glmnet model. The lasso approach requires finding the best value of lambda, which I found by both plotting and assessing the model itself. The lasso model chooses 6 more predictors than the stepwise approach. Yet, the predictors chosen tend to show less significance than the original 6, and removing them would just leave you back at where stepwise was with an AIC of 640, 8 points lower (better) than the current LASSO results at an AIC of 648. Comparing this to our previous stepwise and base models, it would appear that the lasso approach is only slightly better than the base model, while being 8 points in AIC worse than the stepwise model.

Lasso Model Results:

1. Preds: All but P02, LF, and Time. 12 predictors, seems high!
2. AIC = 648
3. R-Squared = 0.615
4. Alpha = 1, Lambda = 6


```

set.seed(1)
#scale all data except categorical So and response
uscrime.scaled = as.data.frame(scale(uscrime[ , -c(2,16)]))
uscrime.scaled <- cbind(uscrime[,2], uscrime.scaled, uscrime$Crime)
# re-insert column for So and Crime
colnames(uscrime.scaled)[1] <- "So"
colnames(uscrime.scaled)[16] <- "Crime"

```

```

scaled.matrix = data.matrix(uscrime.scaled[,-16])
crime.matrix = data.matrix(uscrime.scaled$Crime)

```

```

lasso.model = cv.glmnet( x = scaled.matrix,
                        y = crime.matrix,
                        alpha = 1,
                        nfolds = 6,
                        nlambda = 20 ,
                        type.measure = "mse",
                        family = "gaussian")

```

```
lasso.model
```

```

##
## Call: cv.glmnet(x = scaled.matrix, y = crime.matrix, type.measure = "mse",      nfolds = 6,
##               alpha = 1, nlambda = 20, family = "gaussian")
##
## Measure: Mean-Squared Error
##
##      Lambda Measure      SE Nonzero
## min    5.44    79430 15855      12
## 1se   23.31    95141 20106       9

```


```
summary(lasso.model)
```

```

##      Length Class  Mode
## lambda    20     -none- numeric
## cvm       20     -none- numeric
## cvsd      20     -none- numeric
## cvup      20     -none- numeric
## cvlo      20     -none- numeric
## nzero     20     -none- numeric
## call       8     -none- call
## name       1     -none- character
## glmnet.fit 12     elnet  list
## lambda.min 1     -none- numeric
## lambda.1se 1     -none- numeric

```

```
#plot mean squared error vs lambda. Helps determine which lambda is the best to use.
#appears to be close to a value of 6.
plot(lasso.model)
```



```
#Run again to test those coefficients - a few that were low or negative values
lasso.improv <- lm(Crime~So+M+Ed+Po1+M.F+U2+Ineq, data = uscrime.scaled)
summary(lasso.improv)
```

```
##
## Call:
## lm(formula = Crime ~ So + M + Ed + Po1 + M.F + U2 + Ineq, data = uscrime.scaled)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -441.1  -96.8  -23.6   106.8   623.0
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    911.4      49.4   18.44 < 2e-16 ***
## So             -18.5     111.6   -0.17  0.86917
## M              122.5      48.5    2.53  0.01570 *
## Ed             200.0      65.5    3.05  0.00408 **
## Po1            372.3      44.1    8.44  2.5e-10 ***
## M.F            26.7       40.2    0.66  0.51002
## U2             70.0       39.0    1.80  0.08022 .
## Ineq           250.8      66.6    3.76  0.00055 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 217 on 39 degrees of freedom
## Multiple R-squared:  0.733, Adjusted R-squared:  0.685
## F-statistic: 15.3 on 7 and 39 DF, p-value: 1.93e-09
```

```
# cross validate using 4 folds. MS: 56355
cv.lm.lasso <- cv.lm(lasso.improv, data = uscrime.scaled, m = 4, plotit = FALSE)
```

```

## Analysis of Variance Table
##
## Response: Crime
##      Df Sum Sq Mean Sq F value Pr(>F)
## So      1   56527    56527   1.20 0.27978
## M       1   13927    13927   0.30 0.58951
## Ed      1  905668   905668  19.25 8.5e-05 ***
## Po1     1 3076033 3076033  65.37 7.3e-10 ***
## M.F     1  209271   209271   4.45 0.04143 *
## U2      1  117573   117573   2.50 0.12202
## Ineq    1  666809   666809  14.17 0.00055 ***
## Residuals 39 1835119    47054
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##
## fold 1
## Observations in test set: 11
##      2      9 14      16      20      22      26      38 41      44 47
## Predicted 1355 750.3 711 977.5 1196.9 796 1926 576.1 716 1136.3 1061
## cvpred    1335 762.5 728 1022.5 1205.9 844 1858 603.1 674 1119.8 1047
## Crime     1635 856.0 664 946.0 1225.0 439 1993 566.0 880 1030.0 849
## CV residual 300 93.5 -64 -76.5 19.1 -405 135 -37.1 206 -89.8 -198
##
## Sum of squares = 381988      Mean square = 34726      n = 11
##
## fold 2
## Observations in test set: 12
##      1      3      6      11      19      25      28      29      30      33      35 39
## Predicted 863.5 459 745 1051 1136 643.4 1233 1484 719.6 823 720 718
## cvpred    843.8 357 866 1166 1355 613.8 1202 1725 659.8 752 861 685
## Crime     791.0 578 682 1674 750 523.0 1216 1043 696.0 1072 653 826
## CV residual -52.8 221 -184 508 -605 -90.8 14 -682 36.2 320 -208 141
##
## Sum of squares = 1351124      Mean square = 112594      n = 12
##
## fold 3
## Observations in test set: 12
##      4      5      10      12      13      15      17      34      37      40      42
## Predicted 1844 1270.5 804 611 706 857.3 621.5 998.9 915.4 1075.9 492.4
## cvpred    1818 1259.1 818 622 732 889.7 633.8 1005.2 925.9 1089.5 512.4
## Crime     1969 1234.0 705 849 511 798.0 539.0 923.0 831.0 1151.0 542.0
## CV residual 151 -25.1 -113 227 -221 -91.7 -94.8 -82.2 -94.9 61.5 29.6
##
##      45
## Predicted 583
## cvpred    628
## Crime     455
## CV residual -173
##
## Sum of squares = 204236      Mean square = 17020      n = 12
##
## fold 4
## Observations in test set: 12

```

```
##           7      8      18      21      23      24      27      31      32      36      43      46
## Predicted  707 1326 1118 692.3  864 938.32 274.8  467 764.5  981 1020  809
## cvpred    708 1280 1179 671.4  823 959.63 317.5  623 704.3 1002 1078  890
## Crime      963 1555  929 742.0 1216 968.00 342.0  373 754.0 1272  823  508
## CV residual 255  275 -250  70.6  393   8.37  24.5 -250  49.7  270 -255 -382
##
## Sum of squares = 711356      Mean square = 59280      n = 12
##
## Overall (Sum over all 12 folds)
##      ms
## 56355
```

```
#615
R2(56355)
```

```
## [1] 0.615
```

```
#648
AIC(lasso.improv)
```

```
## [1] 648
```

Elastic Net Model

For the elastic net approach, the challenge was to find the best value of alpha as well. To do this, I made a basic for-loop to run through the model for different values of alpha. This resulted in an alpha value of 0.42 That I then used to find the value of lambda by creating another elastic model with the best chosen value of alpha. This value came out to be a lambda of 13. The resulting r-squared ended up being 0.42 which should cause for some suspicion. Upon further digging, I have come to believe that the insignificant p-values are the reason behind this low r-squared score. If one were to include only the significant values based off of the p-value significance, you would likely end up with a r-squared closer to what was achieved in lasso or stepwise.

In summary, the best approach for this data set is stepwise! It is also likely that the elastic/lasso models are overfit due to the number of predictors that was chosen by for each approach compared to the data set size. Despite this, it may be that the Lasso and Elastic net methods perform better when predicting a new point/city despite the current model results. One would be able to verify this by splitting the data into test and training sets and then using it to predict on a new point as we have done in past lessons, but this is beyond the scope of this homework. Additionally, I wanted to use as much data as possible for the variable selection process due to the relatively small dataset we are working with.

Elastic Net Results:

1. Preds: So+M+Ed+Po1+M.F+Pop+NW+U1+U2+Wealth+Ineq+Prob. Everything but Time. Far more than the previous approaches
2. AIC = 645
3. R-squared = 0.477
4. Alpha = 0.42 and Lambda = 13

```

set.seed(1)
r.squared <- c()

for (i in 0:100) {

  elastic.model <- cv.glmnet( x = scaled.matrix,
                             y = crime.matrix,
                             alpha = i/100,
                             nfolds = 6 ,
                             type.measure = "mse",
                             family = "gaussian")

  fit.select = which(elastic.model$glmnet.fit$lambda == elastic.model$lambda.min)
  #Calc for r-squared
  r.squared <- cbind(r.squared, elastic.model$glmnet.fit$dev.ratio[fit.select])

}
r.squared

```

```

##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12]
## [1,] 0.774 0.708 0.766 0.76 0.748 0.769 0.755 0.761 0.771 0.745 0.765 0.748
##      [,13] [,14] [,15] [,16] [,17] [,18] [,19] [,20] [,21] [,22] [,23] [,24]
## [1,] 0.729 0.766 0.747 0.766 0.764 0.766 0.767 0.769 0.75 0.785 0.766 0.759
##      [,25] [,26] [,27] [,28] [,29] [,30] [,31] [,32] [,33] [,34] [,35] [,36]
## [1,] 0.779 0.741 0.758 0.767 0.756 0.765 0.792 0.776 0.774 0.772 0.765 0.731
##      [,37] [,38] [,39] [,40] [,41] [,42] [,43] [,44] [,45] [,46] [,47] [,48]
## [1,] 0.779 0.74 0.771 0.785 0.755 0.689 0.803 0.766 0.783 0.789 0.764 0.795
##      [,49] [,50] [,51] [,52] [,53] [,54] [,55] [,56] [,57] [,58] [,59] [,60]
## [1,] 0.781 0.749 0.794 0.719 0.783 0.79 0.7 0.721 0.781 0.788 0.764 0.769
##      [,61] [,62] [,63] [,64] [,65] [,66] [,67] [,68] [,69] [,70] [,71] [,72]
## [1,] 0.774 0.716 0.78 0.778 0.788 0.767 0.75 0.779 0.772 0.793 0.752 0.721
##      [,73] [,74] [,75] [,76] [,77] [,78] [,79] [,80] [,81] [,82] [,83] [,84]
## [1,] 0.73 0.759 0.712 0.789 0.796 0.748 0.789 0.792 0.734 0.792 0.789 0.793
##      [,85] [,86] [,87] [,88] [,89] [,90] [,91] [,92] [,93] [,94] [,95] [,96]
## [1,] 0.792 0.749 0.762 0.762 0.775 0.775 0.767 0.763 0.793 0.664 0.784 0.763
##      [,97] [,98] [,99] [,100] [,101]
## [1,] 0.787 0.751 0.682 0.768 0.776

```

```

#now to determine the best value for alpha
alpha.best <- (which.max(r.squared)-1) /100


alpha.best

```

```
## [1] 0.42
```

```
elastic.model <- cv.glmnet( x = scaled.matrix,
  y = crime.matrix,
  alpha = alpha.best,
  nfolds = 6 ,
  nlambda = 20,
  type.measure = "mse",
  family = "gaussian")
```

```
plot(elastic.model)
```



#13 seems to be the best value of lambda according to the model/plot
 elastic.model

```
##
## Call: cv.glmnet(x = scaled.matrix, y = crime.matrix, type.measure = "mse",      nfolds = 6,
##               alpha = alpha.best, nlambda = 20, family = "gaussian")
##
## Measure: Mean-Squared Error
##
##      Lambda Measure   SE Nonzero
## min   13.0   67336 8853      14
## 1se   34.2   71561 7290      12
```

#selects 13 compared, much more than the previous lasso and stepwise approach
coefficients(elastic.model, s=elastic.model\$lambda.min)

```
## 16 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept) 8.88e+02
## So          5.09e+01
## M           9.16e+01
## Ed          1.41e+02
## Po1         2.27e+02
## Po2         5.66e+01
## LF          3.06e+00
## M.F         6.35e+01
## Pop        -6.68e-04
## NW          1.72e+01
## U1         -5.46e+01
## U2          9.13e+01
## Wealth      3.00e+01
## Ineq        1.95e+02
## Prob       -8.72e+01
## Time        .
```

```
elastic.model.improv <- lm(Crime~So+M+Ed+Po1+M.F+Pop+NW+U1+U2+Wealth+Ineq+Prob, data = uscrime)
summary(elastic.model.improv)
```



```
##
## Call:
## lm(formula = Crime ~ So + M + Ed + Po1 + M.F + Pop + NW + U1 +
##      U2 + Wealth + Ineq + Prob, data = uscrime)
##
## Residuals:
##      Min        1Q    Median        3Q       Max
## -434.2  -107.0   18.6   115.9   470.3
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -6.39e+03  1.41e+03  -4.52  7.1e-05 ***
## So           2.29e+01  1.25e+02   0.18  0.8562
## M            8.97e+01  3.93e+01   2.28  0.0288 *
## Ed           1.75e+02  5.63e+01   3.11  0.0038 **
## Po1          9.87e+01  2.19e+01   4.51  7.3e-05 ***
## M.F          1.66e+01  1.63e+01   1.02  0.3166
## Pop         -8.74e-01  1.20e+00  -0.73  0.4711
## NW           1.86e+00  5.61e+00   0.33  0.7419
## U1          -4.98e+03  3.64e+03  -1.37  0.1807
## U2           1.67e+02  7.91e+01   2.11  0.0424 *
## Wealth       8.63e-02  9.90e-02   0.87  0.3893
## Ineq         7.16e+01  2.14e+01   3.35  0.0020 **
## Prob        -4.08e+03  1.81e+03  -2.26  0.0307 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 203 on 34 degrees of freedom
## Multiple R-squared:  0.797, Adjusted R-squared:  0.726
## F-statistic: 11.1 on 12 and 34 DF, p-value: 1.52e-08
```

```
# cross validate using 4 folds
cv.lm.elastic <- cv.lm(elastic.model.improv, data = uscrime, m = 4, plotit = FALSE)
```

```

## Analysis of Variance Table
##
## Response: Crime
##      Df  Sum Sq Mean Sq F value Pr(>F)
## So      1    56527    56527   1.38 0.24880
## M       1    13927    13927   0.34 0.56412
## Ed      1   905668   905668  22.06 4.2e-05 ***
## Po1     1  3076033  3076033  74.92 4.1e-10 ***
## M.F     1   209271   209271   5.10 0.03050 *
## Pop     1    66764    66764   1.63 0.21088
## NW      1    15839    15839   0.39 0.53866
## U1      1      17      17     0.00 0.98400
## U2      1   299228   299228   7.29 0.01074 *
## Wealth  1    43104    43104   1.05 0.31277
## Ineq    1   589797   589797  14.37 0.00059 ***
## Prob    1   208855   208855   5.09 0.03065 *
## Residuals 34 1395898    41056
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##
## fold 1
## Observations in test set: 11
##      2  9  14  16  20  22  26  38  41  44  47
## Predicted 1442 705 770 991 1224 659 1942.9 539.0 785 1189 1063
## cvpred    1517 711 843 1088 1343 829 2061.4 594.2 754 1235 1188
## Crime     1635 856 664 946 1225 439 1993.0 566.0 880 1030 849
## CV residual 118 145 -179 -142 -118 -390 -68.4 -28.2 126 -205 -339
##
## Sum of squares = 431883    Mean square = 39262    n = 11
##
## fold 2
## Observations in test set: 12
##      1  3  6  11  19  25  28  29  30  33  35  39
## Predicted 762.9 345 790 1204 1184 592.6 1236.1 1310 677 888 698 790.8
## cvpred    718.9 316 892 1207 1464 545.6 1167.2 1648 588 835 881 779.3
## Crime     791.0 578 682 1674 750 523.0 1216.0 1043 696 1072 653 826.0
## CV residual 72.1 262 -210 467 -714 -22.6 48.8 -605 108 237 -228 46.7
##
## Sum of squares = 1337662    Mean square = 111472    n = 12
##
## fold 3
## Observations in test set: 12
##      4  5  10  12  13  15  17  34  37  40  42  45
## Predicted 1796 1145 753.3 730.5 746 919 393 996.6 1026 1101.6 296 629
## cvpred    1646 970 658.8 757.7 761 1107 165 949.4 1095 1138.2 103 808
## Crime     1969 1234 705.0 849.0 511 798 539 923.0 831 1151.0 542 455
## CV residual 323 264 46.2 91.3 -250 -309 374 -26.4 -264 12.8 439 -353
##
## Sum of squares = 870267    Mean square = 72522    n = 12
##
## fold 4
## Observations in test set: 12

```

```
##          7      8      18      21      23      24      27      31      32      36      43      46
## Predicted  889.6 1354   863 789.3   954 855 306.5  412   774 1125 1107   789
## cvpred    876.5 1061  1193 771.1   921 830 312.4  517   881   895 1245   923
## Crime      963.0 1555   929 742.0  1216 968 342.0  373   754 1272   823   508
## CV residual 86.5   494 -264 -29.1   295 138   29.6 -144 -127   377 -422 -415
##
## Sum of squares = 957736      Mean square = 79811      n = 12
##
## Overall (Sum over all 12 folds)
##      ms
## 76544
```

```
AIC(elastic.model.improv)
```

```
## [1] 645
```

```
#0.477 R-SQUARED
R2(76544)
```

```
## [1] 0.477
```